


HPC Lab

Session 4: Profiler

Carsten Uphoff, Chaulio Ferreira, Michael Bader
TUM – SCCS

4 December 2017



TUM Uhrenturm

Profiling

Profiling allows you to learn where your program spent its time [...]. This information can show you which pieces of your program are slower than you expected, and might be candidates for rewriting to make your program execute faster.

Source: <https://sourceware.org/binutils/docs/gprof/Introduction.html>

GNU gprof

- Compile the program with profiling enabled:
`icpc -pg program.c -o program`
- Execute the program and generate profile data:
`./program`
(will generate a file `gmon.out`)
- Analyze the data with `gprof`:
`gprof program`

GNU gprof - Flat profile

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
21.75	0.05	0.05	5050000	0.01	0.01	solver::vec2<float>::vec2(float)
13.05	0.08	0.03	5000	6.00	6.00	WavePropagation::updateUnknown
8.70	0.10	0.02	1010000	0.02	0.03	solver::FWave<float>::f(solver
8.70	0.12	0.02	1010000	0.02	0.02	solver::matrix2x2<float>::matr
...						

GNU gprof - Call tree

Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 4.35% of 0.23 seconds

index	% time	self	children	called	name
...					
		0.02	0.00	5001/5001	main [1]
[12]	8.7	0.02	0.00	5001	writer::VtkWriter::write(float, float const*, float c
		0.00	0.00	5001/5001	writer::VtkWriter::generateFileName() [33]
		0.00	0.00	5001/5002	std::operator (std::_Ios_Openmode, std::_Ios_Oper
...					

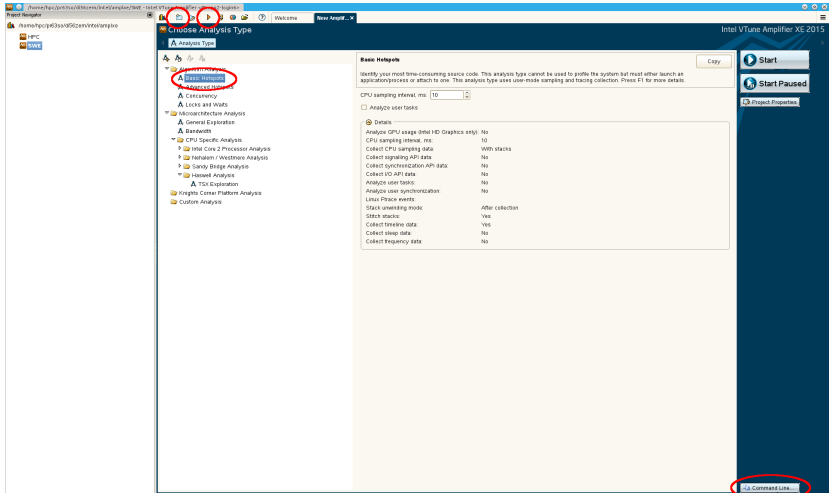
Intel VTune Amplifier XE


On the cluster:

- `module load amplifier_xe`
- GUI: `amplxe-gui`
- Command line tool: `amplxe-cl`

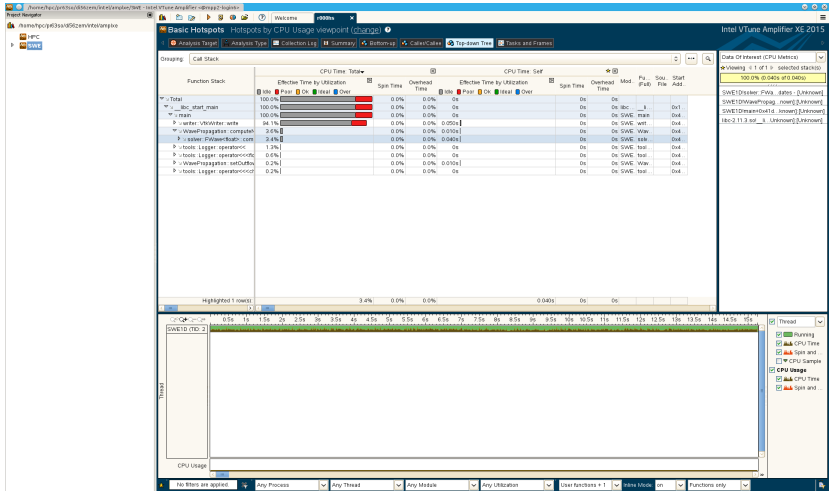
Example: `amplxe-cl -collect hotspots ./program`

Intel VTune Amplifier XE





Intel VTune Amplifier XE



Scalasca

Open source project:
Forschungszentrum Jülich,
Technische Universität Darmstadt,
German Research School for Simulation Sciences

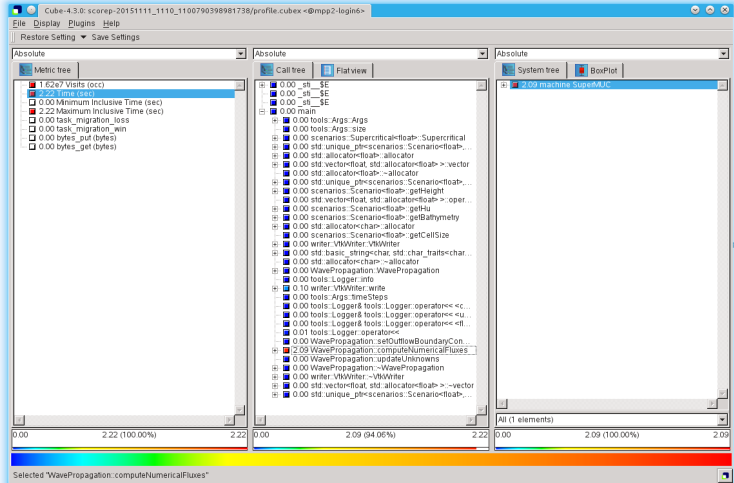
On the Linux Cluster:

- `module load scalasca`
- Also loads:
 - Score-P (Code instrumentation)
 - Cube (Visualisation)

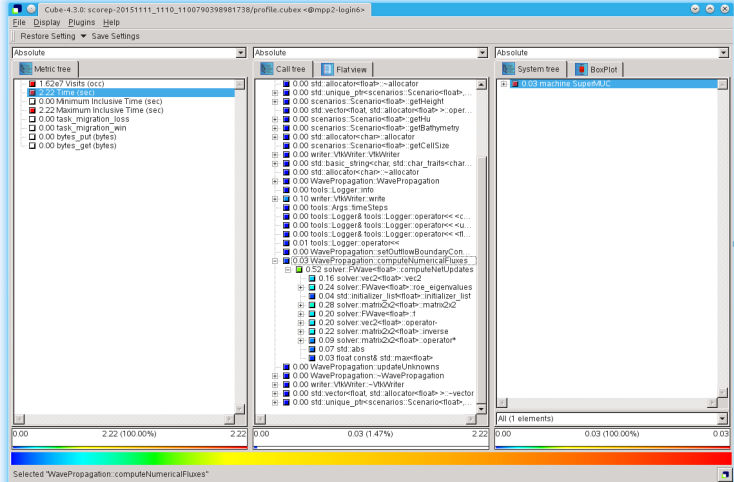
Scalasca – instrumentation

- `scalasca -instrument [options] compiler ...`
Installation on the Linux Cluster only works with the Intel compiler
Custom installation for GCC possible
- Options:
 - `--mpp=mpi`
 - `--thread=omp`
 - `--nocompiler`
 - `--user`
 - ...
- Instrumented code generates a folder `scorep-*`

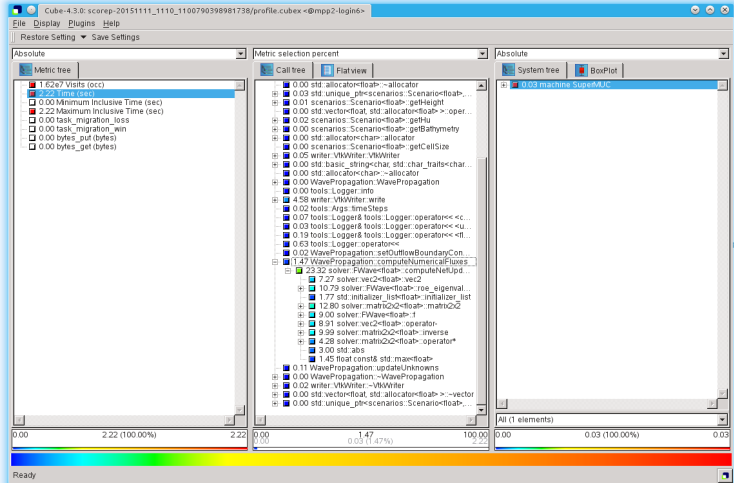
Cube



Cube



Cube



Score-P – Manual Instrumentation

- Option: --user
- Functions:

```
1 #include <scorep/SCOREP_User.h>
2
3 void foo(x) {
4     SCOREP_USER_REGION( "foo", SCOREP_USER_REGION_TYPE_FUNCTION )
5     // Do something
6 }
```

- Regions:

```
1 #include <scorep/SCOREP_User.h>
2 void foo() {
3     SCOREP_USER_REGION_DEFINE( handle )
4     // Do something
5     SCOREP_USER_REGION_BEGIN( handle, "region", SCOREP_USER_REGION_TYPE_COMMON )
6     // Do something else
7     SCOREP_USER_REGION_END( handle )
8     // Do more
9 }
```

Score-P – Parameter-Based Profiling

```
1 #include <scorep/SCOREP_User.h>
2
3 void foo(int64_t myint)
4 {
5     SCOREP_USER_REGION_DEFINE( handle )
6     SCOREP_USER_REGION_BEGIN( handle, "foo", SCOREP_USER_REGION_TYPE_COMMON )
7     SCOREP_USER_PARAMETER_INT64( "myint", myint )
8
9     // do something
10
11     SCOREP_USER_REGION_END( handle )
12 }
```


Hardware Performance Counters

- Hardware counters are special registers
- Count events, e.g.
 - total instructions
 - cache misses
 - branch misses
 - ...
- Automatically incremented by the hardware
 - no instrumentation required
 - minimal overhead
- Likwid-perfctr
 - Linux tool/kernel module to read hardware counters
 - Developed by Regionales Rechenzentrum Erlangen (RRZE)

Likwid-perfctr

```

1 (...)
2 -----
3 DGEMM with options:
4 Time/FLOPS with timeofday(): 1.020164e+10
5 -----
6 Group 1: MEM
7 +-----+
8 |          Event          | Counter | Core 1 |
9 +-----+
10 | INSTR_RETIRED_ANY | FIXC0 | 68146405143 |
11 | CPU_CLK_UNHALTED_CORE | FIXC1 | 52138809008 |
12 | CPU_CLK_UNHALTED_REF | FIXC2 | 41113421180 |
13 (...)
14 +-----+
15 |          Metric          | Core 1 |
16 +-----+
17 (...)
18 | Memory read bandwidth [MBytes/s] | 1.014995e+04 |
19 | Memory read data volume [GBytes] | 1.621715e+02 |
20 | Memory write bandwidth [MBytes/s] | 4.532920e+01 |
21 | Memory writo data volume [GBytes] | 0.72425024 |
22 | Memory bandwidth [MBytes/s] | 1.019528e+04 |
23 | Memory data volume [GBytes] | 1.628957e+02 |
24 +-----+

```

Likwid-perfctr

```

1 (...)
2 -----
3 DGEMM with options: blocked packed vector
4 Time/FLOPS with timeofday(): 2.770782e+10
5 -----
6 Group 1: MEM
7 +-----+
8 |          Event          | Counter | Core 1 |
9 +-----+
10 | INSTR_RETIRED_ANY | FIXC0 | 38479616403 |
11 | CPU_CLK_UNHALTED_CORE | FIXC1 | 19245252158 |
12 | CPU_CLK_UNHALTED_REF | FIXC2 | 15179510632 |
13
14 (...)
15 +-----+
16 |          Metric          | Core 1 |
17 +-----+
18 (...)
19 | Memory read bandwidth [MBytes/s] | 5.727931e+03 |
20 | Memory read data volume [GBytes] | 34.298501696 |
21 | Memory write bandwidth [MBytes/s] | 3.039028e+02 |
22 | Memory write data volume [GBytes] | 1.81975136 |
23 | Memory bandwidth [MBytes/s] | 6.031834e+03 |
24 | Memory data volume [GBytes] | 36.118253056 |
25 +-----+

```

Things to keep in mind

- Measurements almost always create overhead.
- Check if a metric conforms with your performance model.
- Do not trust measurements blindly.